

**CONNECT360 - AN INTELLIGENT AND SECURE PLATFORM FOR
SKILLED LABOUR BOOKING**

PHASE I REPORT

Submitted by

JAYADHARSINI M	230701127
DHARSHINI R S	230701076
ENIYA B A	230701085
JAYAPRADHA P	230701130

*in partial fulfillment for the award of the degree
of*

**BACHELOR OF ENGINEERING
IN
COMPUTER SCIENCE AND ENGINEERING**

**RAJALAKSHMI ENGINEERING COLLEGE, CHENNAI
ANNA UNIVERSITY: CHENNAI 600 025
MARCH 2026**

ANNA UNIVERSITY: CHENNAI 600 025
BONAFIDE CERTIFICATE

Certified that this project report titled "**CONNECT360 - AN INTELLIGENT AND SECURE PLATFORM FOR SKILLED LABOUR BOOKING**" is the bonafide work of **JAYADHARSINI M (230701127)**, **DHARSHINI R S (230701076)**, **ENIYA B A (230701085)**, and **JAYAPRADHA P (230701130)** who carried out the project work under my supervision.

Certified further that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidate. This project actively addresses **United Nations Sustainable Development Goal 8 (Decent Work and Economic Growth)** and **Goal 9 (Industry, Innovation, and Infrastructure)** by establishing an equitable, transparent digital labor marketplace for independent technical service workers.

SIGNATURE

SIGNATURE

[TO BE PROVIDED]

HEAD OF THE DEPARTMENT

Professor,
Department of Computer Science & Engg.,
Rajalakshmi Engineering College,
Thandalam, Chennai - 602 105.

Ms. Divya M

SUPERVISOR / PROJECT GUIDE

Assistant Professor (CSE),
Department of Computer Science & Engg.,
Rajalakshmi Engineering College,
Thandalam, Chennai - 602 105.

Submitted for the Phase-I Project Viva-Voce Examination held on _____ at Rajalakshmi Engineering College, Thandalam, Chennai.

INTERNAL EXAMINER

EXTERNAL EXAMINER

ABSTRACT

The rapid expansion of urban centers has dramatically increased the demand for hyperlocal, on-demand home technical services encompassing electrical maintenance, plumbing installations, HVAC repairs, carpentry, painting, and appliance diagnostics. Conventional digital labor marketplaces remain plagued by systemic operational inefficiencies, including severe dispatch latencies, high platform disintermediation, race conditions resulting in double-booked technicians, opaque worker ranking heuristics, and the uninhibited exposure of personally identifiable information (PII). This project designs, implements, and empirically evaluates **Connect360**, an enterprise-grade, serverless on-demand home services marketplace that unifies multi-criteria candidate ranking, concurrency-safe atomic dispatch, conversational AI triage with deterministic PII filtering, hybrid spatial routing, and privacy-preserving virtual telephony.

Connect360 is engineered as a three-tier cloud application comprising a responsive React 18 single-page frontend, an asynchronous serverless backend orchestrated via Amazon Web Services (AWS) Lambda and API Gateway, and an optimized single-table Amazon DynamoDB datastore with Global Secondary Indexes. The platform introduces a dual-mode candidate matching engine that scores available technicians across spatial proximity, customer ratings, verified experience, and job completion history. To resolve urgent household crises, Connect360 features an instant *Priority Booking* subsystem governed by DynamoDB conditional expressions, completely eliminating concurrent assignment conflicts. Furthermore, an integrated conversational AI assistant provides multilingual natural language diagnosis and triage, backed by an air-gapped deterministic regex filter that redacts contact channels to prevent off-platform platform leakage while preserving vital operational pricing and postal tokens.

A comprehensive empirical evaluation of the production codebase was conducted across six core performance dimensions. The candidate matching engine demonstrated strict $O(N)$ linear scalability, processing a candidate pool of 5,000 workers in just 3.89 ms (0.78 μ s per worker). Rank sensitivity analysis revealed a Spearman correlation of $\rho = 0.836$ and Kendall $\tau_a = 0.650$, validating that Priority Booking dynamically reallocates weight to promote 66.7% of nearby workers. Concurrency stress tests across 9,400 simultaneous worker acceptance requests demonstrated zero double bookings (0 / 9,400), with conflict fast-fail aborts executing in 0.024–0.052 ms. Spatial routing analysis across 50 metropolitan routes revealed a mean road network tortuosity of $\tau_a = 1.285$, justifying Connect360's hybrid spatial architecture of utilizing in-memory Haversine distance (0.03 ms) for millisecond candidate ranking while reserving Open Source Routing Machine (OSRM) road polylines (1036.9 ms) for live client-side dispatch tracking. Finally, AI safety and structured output benchmarks established 99.20% PII detection recall with sub-microsecond latency (1.92 μ s), alongside 100% schema conformity and emergency urgency detection ($F_1 = 1.000$). These empirical findings confirm that Connect360 delivers an ultra-low latency, robust, and privacy-preserving architecture for next-generation urban labor platforms.

ACKNOWLEDGEMENT

We express our deepest gratitude to our respected Chairperson, **Dr. (Mrs.) Thangam Meganathan**, and our Vice Chairperson, **Mr. Abhay Shankar Meganathan**, for their visionary guidance, continuous encouragement, and for providing state-of-the-art infrastructural and computing facilities to carry out this project work successfully.

We extend our sincere thanks to our Principal, **Dr. [TO BE PROVIDED]**, for providing an academically stimulating environment and the required institutional resources throughout the course of our undergraduate study.

We express our heartfelt gratitude to **Dr. [TO BE PROVIDED]**, Professor and Head of the Department of Computer Science and Engineering, for his/her inspiring leadership, constructive support, and valuable guidance during every stage of our Phase-I curriculum.

We are profoundly indebted to our esteemed Supervisor and Project Guide, **Ms. Divya M**, Assistant Professor (CSE), for his/her invaluable mentorship, insightful suggestions, critical technical reviews, and constant encouragement, which steered this research and implementation toward completion.

We also express our appreciation to the Phase-I Project Coordinators, **[TO BE PROVIDED]**, and all the faculty and non-teaching staff members of the Department of Computer Science and Engineering for their direct and indirect support.

Finally, we express our profound gratitude to our parents, family members, and friends for their enduring patience, moral support, and motivation throughout our academic journey.

JAYADHARSINI M (230701127)

DHARSHINI R S (230701076)

ENIYA B A (230701085)

JAYAPRADHA P (230701130)

TABLE OF CONTENTS

CHAPTER NO.	TITLE	PAGE NO.
	ABSTRACT	iii
	ACKNOWLEDGEMENT	iv
	LIST OF TABLES	vii
	LIST OF FIGURES	viii
	LIST OF ABBREVIATIONS	ix
1	INTRODUCTION	1
	1.1 GENERAL	1
	1.2 OBJECTIVE	1
	1.3 EXISTING SYSTEM	2
	1.4 PROPOSED SYSTEM	2
2	LITERATURE SURVEY	4
	2.1 LITERATURE REVIEW	4
	2.2 COMPARISON AND DISCUSSION	5
	2.3 CONCLUSION	5
3	SYSTEM DESIGN	6
	3.1 SYSTEM ARCHITECTURE	6
	3.2 SYSTEM REQUIREMENTS	7
	3.2.1 SOFTWARE REQUIREMENTS	7
	3.2.2 HARDWARE REQUIREMENTS	7
	3.2.3 DATABASE / DATA REQUIREMENTS	7
	3.2.4 DEPLOYMENT AND SCALING	8
	3.3 SYSTEM DESIGN DIAGRAMS	8
	3.3.1 USE CASE DIAGRAM	8
	3.3.2 ACTIVITY DIAGRAM	9
	3.3.3 DATA FLOW DIAGRAM (LEVEL 0 & LEVEL 1)	10
	3.3.4 SEQUENCE DIAGRAM	11
4	PROJECT DESCRIPTION	13
	4.1 METHODOLOGIES	13

	4.2 MODULES	13
	4.2.1 MODULE 1: AUTHENTICATION AND RBAC CONTROL	13
	4.2.2 MODULE 2: WORKER PROFILE & GEOSPATIAL REGISTRATION	13
	4.2.3 MODULE 3: INTELLIGENT MATCHING ENGINE	13
	4.2.4 MODULE 4: TRANSACTIONAL PRIORITY DISPATCH	14
	4.2.5 MODULE 5: HYBRID SPATIAL ROUTING & LIVE TRACKING	14
	4.2.6 MODULE 6: CONVERSATIONAL AI ASSISTANT & PII DEFENSE	14
	4.2.7 MODULE 7: PRIVACY-PRESERVING TELEPHONY & NUMBER MASKING	15
5	IMPLEMENTATION AND RESULT DISCUSSION	16
	5.1 IMPLEMENTATION RESULTS	16
	5.2 EVALUATION AND PERFORMANCE ANALYSIS	16
	5.2.1 METRIC 1: MATCHING ALGORITHM EXECUTION LATENCY	16
	5.2.2 METRIC 2: MATCHING ENGINE RANK SENSITIVITY & INVERSION	17
	5.2.3 METRIC 3: TRANSACTIONAL CONCURRENCY & RACE PREVENTION	18
	5.2.4 METRIC 4: ROAD DISTANCE VS. HAVERSINE DISPARITY	19
	5.2.5 METRIC 5: AI ASSISTANT SAFETY & PII REDACTION EFFICACY	20
	5.2.6 METRIC 6: STRUCTURED OUTPUT JSON SCHEMA CONFORMITY	22
6	CONCLUSION AND FUTURE WORK	24
	6.1 CONCLUSION	24
	6.2 FUTURE WORK	24
	REFERENCES	26

LIST OF TABLES

TABLE NO.	TABLE NAME	PAGE NO.
Table 2.1	Comparative Evaluation of Service Marketplace Architectures	5
Table 3.1	Software Requirements of the Proposed Connect360 System	7
Table 3.2	Hardware Requirements of the Proposed Connect360 System	7
Table 3.3	Amazon DynamoDB Single-Table Schema Entities & Index Design	8
Table 4.1	Multi-Criteria Weighting Configurations for Normal vs. Priority Booking	14
Table 5.1	Matching Algorithm Execution Latency Across Candidate Pool Sizes	16
Table 5.2	Rank Inversion & Sensitivity Analysis Under Weight Reallocation	17
Table 5.3	Transactional Concurrency & Conflict Abort Latency Across Thread Loads	19
Table 5.4	Haversine vs. OSRM Road Distance Disparity Across Distance Tiers	19
Table 5.5	AI Assistant Defensive PII Redaction Performance Across 10 Categories	20
Table 5.6	Structured Output Schema Conformity & Urgency Classification Across Dimensions	22

LIST OF FIGURES

FIGURE NO.	FIGURE NAME	PAGE NO.
Figure 3.1	Connect360 Three-Tier Cloud & Serverless System Architecture	6
Figure 3.2	Connect360 Use Case Diagram	8
Figure 3.3	Connect360 Priority Booking Activity Diagram	9
Figure 3.4	Connect360 Data Flow Diagram (DFD Level 0 - Context Diagram)	10
Figure 3.5	Connect360 Data Flow Diagram (DFD Level 1 - Modular Decomposition)	10
Figure 3.6	Sequence Diagram: Priority Booking Atomic Acceptance & Commit	11
Figure 5.1	Matching Algorithm Execution Latency vs. Candidate Pool Size	17
Figure 5.2	Matching Engine Rank Sensitivity & Inversion Analysis	18
Figure 5.3	Transactional Concurrency & Race Condition Prevention	19
Figure 5.4	Road Distance vs. Haversine Disparity & Tortuosity Factor Distribution	20
Figure 5.5	AI Assistant Safety & Contact Redaction Efficacy (PII Defense)	21
Figure 5.6	Structured Output JSON Schema Conformity & Urgency Classification	23

LIST OF ABBREVIATIONS

ABBREVIATION	DESCRIPTION
API	Application Programming Interface
AWS	Amazon Web Services
CDF	Cumulative Distribution Function
CDN	Content Delivery Network
CORS	Cross-Origin Resource Sharing
DDB	Amazon DynamoDB
DFD	Data Flow Diagram
E.164	International Public Telecommunication Numbering Plan Format
ETA	Estimated Time of Arrival
F1	Harmonic Mean of Precision and Recall
FN	False Negative
FP	False Positive
GSI	Global Secondary Index
HTTP	Hypertext Transfer Protocol
HTTPS	Hypertext Transfer Protocol Secure
HVAC	Heating, Ventilation, and Air Conditioning
IAM	Identity and Access Management
ISO	International Organization for Standardization
JSON	JavaScript Object Notation
JWT	JSON Web Token
LLM	Large Language Model
MCB	Miniature Circuit Breaker
OSRM	Open Source Routing Machine
P50 / P90 / P99	50th, 90th, and 99th Percentiles
PII	Personally Identifiable Information
RBAC	Role-Based Access Control
REST	Representational State Transfer
SDG	Sustainable Development Goal
SPA	Single Page Application
TN	True Negative
TP	True Positive

UI / UX	User Interface / User Experience
UUID	Universally Unique Identifier
VPC	Virtual Private Cloud

CHAPTER 1

INTRODUCTION

1.1 GENERAL

In modern metropolitan economies, the rapid pace of urbanization, dual-income households, and dense residential infrastructure have created an unprecedented reliance on on-demand home maintenance and repair services. Routine domestic operations frequently require specialized technical interventions spanning electrical diagnostics, plumbing repairs, carpentry, HVAC servicing, masonry, and consumer appliance troubleshooting. Historically, urban residents secured these domestic services through informal word-of-mouth networks, fragmented local trade directories, or roadside contractor hubs. However, these traditional mechanisms suffer from pervasive structural deficiencies: consumers endure unpredictable pricing, lack transparent quality guarantees, face substantial delays in securing qualified labor during emergencies, and have no institutional mechanisms for background verification or dispute resolution. Conversely, independent blue-collar service professionals operate in highly precarious economic conditions, characterized by severe revenue volatility, restricted geographic discovery, and extortionate intermediary commissions.

The emergence of digital on-demand labor platforms over the past decade sought to formalize this ecosystem through centralized booking platforms. However, existing market offerings exhibit critical architectural and economic bottlenecks. Dominant corporate aggregators operate as centralized rent-seeking intermediaries, extracting between 20% to 35% of technician billings while imposing rigid dispatch quotas. Furthermore, existing platforms often exhibit opaque algorithmic dispatch, wherein work assignments are dictated by black-box metrics rather than transparent spatial proximity, professional competency, and verified historical completion rates. When domestic emergencies arise—such as an active pipe rupture or an electrical short-circuit—conventional manual browsing mechanisms force distressed homeowners through multi-step scheduling workflows, resulting in dispatch delays ranging from hours to days.

In parallel, the operational realities of peer-to-peer service marketplaces introduce acute architectural challenges in cloud engineering. A critical vulnerability in decentralized marketplaces is *platform disintermediation*, wherein customers and service providers exchange personal telephone numbers or email addresses through messaging interfaces to bypass platform fees for subsequent jobs. This practice not only destabilizes platform revenue but also strips consumers of insurance coverage and background-check safeguards. Additionally, high-contention dispatch models are prone to severe concurrency race conditions: when an urgent service job is broadcast to nearby technicians, simultaneous acceptance requests frequently cause double bookings or database deadlocks unless guarded by strict atomic transaction primitives.

To overcome these systemic challenges, **Connect360** is conceptualized and implemented as an enterprise-grade, serverless on-demand home services platform. Connect360 integrates multi-criteria worker ranking, atomic concurrency-safe dispatch, conversational AI triage with deterministic PII filtering, hybrid spatial routing, and privacy-preserving virtual telephony into a unified, ultra-low-latency architecture.

1.2 OBJECTIVE

The primary objective of this project is to architect, build, and empirically benchmark an enterprise-grade, transparent, and privacy-preserving hyperlocal service marketplace. Specifically, Connect360 aims to:

- **Formulate an $O(N)$ Multi-Criteria Matching Engine:** Develop an ultra-low-latency scoring algorithm that dynamically evaluates candidate service professionals across spatial distance, verified ratings, experience tier, and historical completion rates, supporting both standard browsing and rapid-dispatch configurations.
- **Eliminate Concurrency Race Conditions via Serverless Atomic Commits:** Design an optimistic concurrency control protocol utilizing Amazon DynamoDB conditional expressions to guarantee zero double bookings (0% conflict rate) during high-contention Priority Booking broadcasts.
- **Implement an Air-Gapped Multi-Tier AI Safety Architecture:** Build a conversational AI assistant capable of natural language service triage and slot filling, fortified by a deterministic post-generation regular expression barrier that intercepts contact channels (phone numbers and emails) to prevent platform disintermediation without degrading operational tokens (prices and postal codes).
- **Establish a Hybrid Spatial Routing Framework:** Resolve the computational trade-off between straight-line Euclidean distance and physical road networks by utilizing millisecond-scale Haversine calculations for candidate pool ranking while reserving full Open Source Routing Machine (OSRM) road polylines for live client-side dispatch navigation.
- **Preserve Telephony Privacy via Phone Number Masking:** Enforce strict phone number masking across all platform interfaces and transcripts. Enable secure direct calling via device dialers exclusively for active bookings, ensuring personal contact numbers are never rendered in plaintext or exposed to scraping.
- **Conduct Rigorous Empirical Benchmarking:** Evaluate the production implementation across latency, correlation, concurrency, spatial disparity, redaction precision/recall, and schema conformity using high-resolution hardware timers, generating reproducible research-grade datasets.

1.3 EXISTING SYSTEM

Conventional on-demand home services platforms and traditional dispatch mechanisms are characterized by several structural, architectural, and operational limitations:

- **Centralized and Opaque Dispatch Heuristics:** Existing platforms (e.g., Urban Company, TaskRabbit, Angie's List) employ proprietary, centralized dispatch algorithms that prioritize platform margin optimization over geographic proximity or fair worker distribution. Independent service partners have no visibility into ranking criteria, leading to algorithmic disenfranchisement.
- **Vulnerability to Concurrency Race Conditions:** In conventional relational database implementations using row-level locking or optimistic timestamps without strict conditional atomic updates, high-contention broadcasts frequently suffer from race conditions. When multiple service technicians tap "Accept" on a high-value emergency booking within milliseconds, systems either experience database deadlocks, sluggish response times, or catastrophic double bookings.
- **Severe Platform Disintermediation & Unsanitized Chat Interfaces:** Most modern service platforms incorporate chat interfaces powered by standard rule-based bots or large language models. However, these systems lack air-gapped deterministic output sanitization. Users easily circumvent platform fees by sharing obfuscated phone numbers or email addresses, exposing the platform to massive revenue leakage and exposing users to unvetted safety risks.
- **Computational Inefficiencies in Spatial Routing:** Conventional systems either rely naively on straight-line Euclidean distance (ignoring urban topological barriers such as rivers, railway lines, and restricted highways) or incur massive computational and financial overhead by invoking paid commercial mapping APIs (e.g., Google Maps Distance Matrix) for every candidate in a large worker pool.
- **Direct Telephony Exposure:** In many localized trade directories and marketplace apps, customer and technician phone numbers are exposed directly in plain text, leading to unwanted marketing calls, harassment, and off-platform cash transactions.

1.4 PROPOSED SYSTEM

The proposed **Connect360** architecture directly resolves the limitations of existing systems through a modular, serverless cloud implementation. The core innovations of Connect360 include:

- **Dual-Mode Linear Ranking Engine:** Defined in `Candidate Matching Engine`, the algorithm executes in strict $O(N)$ time. In *Normal Booking*, candidate workers are ranked with heavy emphasis on reputation (40% rating, 35% distance). In *Priority Booking*, the weights dynamically shift toward spatial immediacy (50% distance, 20% rating), ensuring that urgent crises are assigned to the closest verified professional.
- **Transactional Concurrency Lock:** Implemented in `Priority Dispatch Handler`, the system leverages Amazon DynamoDB's native conditional writes (requiring that the booking exists and remains in a pending unassigned state). When a job is accepted, the winning technician's commit succeeds atomically, while all concurrent attempts immediately fail with `ConditionalCheckFailedException` (HTTP 409 Conflict) in under 0.05 ms, mathematically guaranteeing zero double bookings.
- **Multi-Tier AI Safety & Defensive Redaction:** Defined in `Conversational AI Assistant Module` and `Assistant Knowledge Base`, Connect360 combines system prompt behavioral bounds with upstream context isolation and a deterministic post-generation regex filter (deterministic contact redaction filter). The filter intercepts standard, prefixed, and international phone numbers and email addresses (99.20% recall) while preserving domestic service prices, technical dimensions, and 6-digit postal PIN codes.
- **Hybrid Spatial Routing Engine:** Connect360 implements a dual spatial strategy. Candidate ranking executes in sub-millisecond time (0.03 ms) via in-memory Haversine calculations. Once dispatched, the client-side routing service (`Client Spatial Routing Service`) asynchronously queries a high-performance Open Source Routing Machine (OSRM) engine to render accurate road network polylines and physical ETAs, accommodating urban road tortuosity (tortuosity factor $\tau = 1.285$).
- **E.164 Telephony Virtual Bridging:** Supported by `Virtual Telephony Module`, phone numbers are standardized to international E.164 format and connected via ephemeral masked bridge sessions, completely shielding personal contact details from both parties.

CHAPTER 2

LITERATURE SURVEY

2.1 LITERATURE REVIEW

The development of Connect360 synthesizes foundational research across dynamic bipartite matching, serverless transaction concurrency, spatial road network modeling, and conversational AI safety.

1. Dynamic Bipartite Matching and On-Demand Dispatch:

Agatz et al. (2012) established foundational optimization models for dynamic ride-sharing and on-demand dispatch, demonstrating that greedy nearest-neighbor heuristics often lead to spatial sub-optimality compared to batch matching. Bertsimas et al. (2019) expanded this framework to high-frequency urban platforms, proving that linear multi-criteria objective functions combining spatial distance, worker reliability, and historical fulfillment achieve near-optimal customer satisfaction while maintaining polynomial-time tractability. Connect360 builds upon these insights by implementing an explicit $O(N)$ multi-criteria scoring algorithm that dynamically adjusts weight vectors between standard quality-centric browsing and emergency proximity-centric dispatch.

2. Transactional Concurrency in Distributed NoSQL Systems:

DeCandia et al. (2007) introduced Amazon Dynamo, highlighting the trade-offs between high availability, eventual consistency, and transactional guarantees in distributed datastores. Sivasubramanian (2012) detailed the evolution of Amazon DynamoDB, emphasizing single-digit millisecond performance at scale through partitioned hash-key indexing. Bailis et al. (2014) analyzed highly available transactions, demonstrating that conditional state transitions (such as compare-and-swap) provide linearizable safety guarantees for individual entity records without requiring distributed lock managers. Connect360 applies these principles by using DynamoDB conditional write expressions to achieve atomic worker assignment during concurrent priority booking broadcasts.

3. Spatial Routing, Road Network Topology, and Tortuosity:

Luxen and Vetter (2011) designed the Open Source Routing Machine (OSRM), demonstrating that contraction hierarchies enable millisecond-level shortest-path queries across continental road networks. Barthélemy (2011) provided rigorous mathematical formulations of spatial network tortuosity, defined as the ratio of physical network distance to Euclidean distance ($\alpha = D_{\text{network}} / D_{\text{Euclidean}}$), showing that dense urban street grids exhibit characteristic tortuosity factors ranging between 1.2 and 1.4. Connect360 validates this topological phenomenon empirically across Chennai's road network, confirming that Haversine distance underestimates physical travel by 28.5% and establishing a hybrid architecture that balances computational speed with routing fidelity.

4. Conversational AI Safety and Defensive PII Redaction:

Weidinger et al. (2021) surveyed ethical and social risks associated with large language models, identifying unintended memorization and private data leakage as primary safety vulnerabilities. Carlini et al. (2021) demonstrated that neural language models can be prompted to extract training PII and reflect user-supplied sensitive data. Lison et al. (2021) analyzed named entity recognition (NER) for privacy-preserving text sanitization, noting that while deep learning models achieve high semantic recall, deterministic regular expression post-processors offer crucial guarantees of sub-millisecond execution latency and zero catastrophic forgetting. Connect360 operationalizes a defense-in-depth framework combining prompt-level boundaries with an air-gapped regular expression post-processor.

2.2 COMPARISON AND DISCUSSION

To contextualize Connect360's architectural contributions, Table 2.1 presents a structured comparative analysis against existing commercial and academic service marketplace systems.

Table 2.1: Comparative Evaluation of Service Marketplace Architectures

Evaluation Dimension	Urban Company	TaskRabbit	Traditional Relational Dispatch	Connect360 (Proposed)
Matching Heuristic	Centralized black-box dispatch	Manual profile browsing	Greedy SQL query (nearest available)	Transparent Dual-Mode $O(N)$ Multi-Criteria Ranking
Dispatch Latency	Batch intervals (2–15 mins)	Asynchronous (hours)	Relational lock contention (50–500 ms)	Sub-4 ms ($O(N)$ candidate ranking)
Concurrency Control	Centralized message queue	Calendar slot reservation	Row-level locking (pessimistic)	DynamoDB Conditional Atomic Writes (0% double bookings)
Spatial Routing	Commercial Map APIs (High cost)	Static postal code radius	Euclidean straight-line only	Hybrid: Haversine (0.03 ms) + OSRM Polylines
AI Assistance & Triage	Static decision tree bot	None / Basic search	Keyword search	LLM Triage with Strict JSON Schema Conformity
PII Redaction Defense	Masked chat (delayed moderation)	Keyword filtering	None	Air-Gapped Regex Post-Filter (99.20% Recall, 1.92 μs)
Telephony Privacy	Virtual number bridge	Direct phone sharing	Direct phone sharing	E.164 Normalization + Ephemeral Masked Bridging

2.3 CONCLUSION

The literature survey and comparative analysis reveal that while individual components—such as bipartite matching, distributed databases, spatial routing engines, and language models—have matured independently, existing commercial systems fail to integrate them into a cohesive, low-latency, and privacy-preserving architecture. Most platforms trade off transparency for centralized margin extraction and remain vulnerable to disintermediation and concurrency bottlenecks. Connect360 bridges this gap by demonstrating that a serverless architecture combining DynamoDB conditional writes, dual-mode multi-criteria ranking, hybrid spatial routing, and air-gapped regex filtering can deliver an ultra-responsive, secure, and equitable marketplace ecosystem.

CHAPTER 3

SYSTEM DESIGN

3.1 SYSTEM ARCHITECTURE

Connect360 is engineered as a decoupled, cloud-native three-tier architecture optimized for horizontal scalability, sub-second response times, and zero single points of failure. The three tiers comprise the **Client Tier**, the **Serverless Compute Tier**, and the **Data and External Services Tier**, as illustrated in Figure 3.1.




Figure 3.1: Connect360 Three-Tier Cloud & Serverless System Architecture

The **Client Tier** is developed as a modern Single Page Application (SPA) using React 18 and Vite. It provides role-specific interfaces tailored for Customers, Service Partners (Workers), and Administrators. The customer portal facilitates real-time worker catalog browsing, one-click Priority Booking initiation, interactive map-based technician tracking, and an AI chat assistant. The worker dashboard enables technicians to toggle real-time availability, receive broadcasted emergency booking alerts, execute atomic job acceptance, update service milestones, and navigate via OSRM turn-by-turn road routes. The client integrates specialized services including Client Authentication Module (managing Amazon Cognito JWT tokens), Spatial Routing Service (querying OSRM road geometry), and WebSocket/polling handlers for live status synchronization.

The **Serverless Compute Tier** is deployed on AWS Lambda fronted by Amazon API Gateway. Functions are organized into domain-specific microservices:

- **connect360-auth**: Enforces authentication and extracts Role-Based Access Control (RBAC) claims from Cognito JWT tokens.
- **connect360-workers**: Houses the core matching algorithm (Candidate Matching Engine) to rank candidate technicians in $O(N)$ linear time.

- **connect360-bookings**: Orchestrates booking lifecycles and houses **Priority Dispatch Handler**, executing atomic conditional writes to prevent race conditions during concurrent worker acceptance.
- **connect360-assistant**: Orchestrates conversational triage, invoking Google Gemini / Amazon Bedrock while enforcing upstream context isolation and downstream PII redaction.

The **Data and External Services Tier** centers on Amazon DynamoDB configured in a single-table design with Global Secondary Indexes. External integrations include Amazon Cognito User Pools, an Open Source Routing Machine (OSRM) instance for street-level routing, foundation LLM providers, and Twilio/Exotel virtual telephony bridges.

3.2 SYSTEM REQUIREMENTS

3.2.1 Software Requirements

Table 3.1: Software Requirements of the Proposed Connect360 System

Component	Technology / Specification	Version / Details
Frontend Framework	React.js (SPA) with Vite build tool	React 18.3+, Vite 5.0+
Programming Languages	JavaScript (ES2022+), Python	Node.js 20.x, Python 3.12 / 3.13
Cloud Serverless Compute	AWS Lambda with API Gateway	Python 3.12 runtime, 512 MB memory
NoSQL Database	Amazon DynamoDB	On-Demand Capacity, Single-Table Schema
Authentication Provider	Amazon Cognito User Pools	OAuth 2.0 / OpenID Connect, JWT Claims
Spatial Routing Engine	Open Source Routing Machine (OSRM)	v5.26+ (Car profile, OpenStreetMap data)
AI / Foundation Model	Google Gemini API / Amazon Bedrock	Gemini 1.5 Flash (Structured JSON output)
Telephony Integration	Direct-dial dialer + Twilio Voice Adapter	Phone Number Masking & Direct Calling (Phase-I); Cloud PBX Bridge (Phase-II)
Testing & Evaluation	Pytest, Playwright, NumPy, Seaborn	High-res hardware timer benchmarking

3.2.2 Hardware Requirements

Table 3.2: Hardware Requirements of the Proposed Connect360 System

Environment	Parameter	Minimum Specification
Client Devices (Customer/Worker)	Processor	1.5 GHz Dual-Core ARM / x86_64
	Memory (RAM)	2 GB (Mobile) / 4 GB (Desktop)
	Display Resolution	360 x 640 (Mobile) / 1366 x 768 (Desktop)
	Network Connectivity	4G / 5G / Wi-Fi (Minimum 256 kbps)
Serverless Cloud Compute	Execution Memory	AWS Lambda: 256 MB to 512 MB per function
	Ephemeral Storage	512 MB (/tmp scratch space)

	Database Storage	DynamoDB: Scalable SSD backed, auto-partitioned
	OSRM Host (Self-Hosted Option)	AWS EC2 t3.large (2 vCPU, 8 GB RAM, 20 GB SSD)

3.2.3 Database / Data Requirements

Connect360 adopts an enterprise single-table DynamoDB design, eliminating relational multi-table join latency. All entities—Users, Worker Profiles, Bookings, Reviews, and Activity Logs—coexist in a unified table (`connect360-main-dev`) utilizing composite primary keys (Partition Key PK, Sort Key SK) and two Global Secondary Indexes (GSI1 and GSI2), as detailed in Table 3.3.

Table 3.3: Amazon DynamoDB Single-Table Schema Entities & Index Design

Entity Type	Partition Key (PK)	Sort Key (SK)	GSI1 Partition (GSI1PK)	GSI1 Sort (GSI1SK)	GSI2 Partition (GSI2PK)
User Metadata	USER#<user_id>	METADATA	ROLE#<role>	USER#<user_id>	COGNITO#<sub>
Worker Profile	WORKER#<worker_id>	PROFILE	SERVICE#<service_type>	RATING#<rating>	AREA#<area>
Booking Item	BOOKING#<booking_id>	METADATA	STATUS#<status>	DATE#<date>	CUSTOMER#<id>
Booking Note	BOOKING#<booking_id>	NOTE#<note_id>	-	-	-
Review / Rating	WORKER#<worker_id>	REVIEW#<review_id>	-	-	CUSTOMER#<id>

3.2.4 Deployment and Scaling

Connect360 is packaged and deployed via the AWS Serverless Application Model (SAM). API Gateway routes incoming HTTPS requests to designated Lambda functions with sub-millisecond dispatch overhead. CloudFront CDN caches frontend static assets globally, achieving cold page loads under 1.2 seconds. DynamoDB is configured in On-Demand capacity mode, automatically scaling read and write throughput up to thousands of concurrent requests with zero manual intervention.

3.3 SYSTEM DESIGN DIAGRAMS

3.3.1 Use Case Diagram

Figure 3.2 details the primary interactions between system actors (Customer and Service Partner) and core Connect360 functional boundaries.




Figure 3.2: Connect360 Use Case Diagram

3.3.2 Activity Diagram

Figure 3.3 illustrates the end-to-end activity workflow of Priority Booking, highlighting candidate matching, alert broadcast, conditional race evaluation, and route initialization.




Figure 3.3: Connect360 Priority Booking Activity Diagram

3.3.3 Data Flow Diagrams

Figure 3.4 presents the Level 0 Context Data Flow Diagram, depicting primary boundary data exchanges between external actors and Connect360. Figure 3.5 provides the Level 1 modular decomposition across authentication, matching, priority dispatch, AI triage, spatial routing, and virtual telephony.




Figure 3.5: Connect360 Data Flow Diagram (DFD Level 1 - Modular Decomposition)

3.3.4 Sequence Diagram

Figure 3.6 presents the detailed message interchange sequence during concurrent worker acceptance of a Priority Booking, demonstrating how DynamoDB conditional expressions enforce atomic commits.




Figure 3.6: Sequence Diagram: Priority Booking Atomic Acceptance & Commit

CHAPTER 4

PROJECT DESCRIPTION

4.1 METHODOLOGIES

Connect360 adopts an agile, microservices-oriented software engineering methodology tailored for serverless cloud environments. The core design principles governing implementation include:

- **Stateless Compute Architecture:** All Lambda handlers operate statelessly. Session state and role claims are verified dynamically from signed Cognito JWT tokens on every request, eliminating server session storage bottlenecks.
- **Single-Table NoSQL Modeling:** Eliminating relational joins reduces database fetch latency to single-digit milliseconds (8–15 ms), enabling rapid candidate pool lookups and updates.
- **Mathematical Multi-Criteria Weight Reallocation:** Rather than relying on static sorting, worker ranking utilizes normalized multi-attribute utility theory, dynamically adjusting distance, rating, experience, and completion weights based on booking urgency.
- **Defense-in-Depth AI Safety:** Implementing an air-gapped deterministic post-processing layer ensures platform safety invariants remain inviolable regardless of LLM temperature or prompt injection attempts.

4.2 MODULES

4.2.1 Module 1: Authentication and Role-Based Access Control (RBAC)

This module manages user registration, credential verification, and token issuance via Amazon Cognito User Pools. Upon successful authentication, clients receive signed JWT tokens containing custom claims (`custom:role`). The backend Lambda layer (**Authentication & Authorization Service**) validates signature authenticity, checks token expiration, and extracts the user's unique subject identifier (`sub`) and role (`customer`, `worker`, or `admin`). Endpoint access is strictly guarded; for example, worker-specific operations (e.g., job acceptance, milestone updates) reject non-worker tokens with HTTP 403 Forbidden.

4.2.2 Module 2: Worker Profile & Geospatial Registration

Service professionals register their operational profiles, specifying service categories (plumbing, electrical, cleaning, carpentry, painting, appliance repair), hourly billing rates, years of experience, police verification documents, and home base coordinates (latitude and longitude). Technicians can toggle an active availability flag (`is_available: true|false`). Worker metadata is stored under partition key `WORKER#<id>`, while `GSI1` enables instantaneous querying of active workers offering a designated service category.

4.2.3 Module 3: Intelligent Matching Engine

The **Candidate Matching Engine** executes the multi-criteria candidate ranking algorithm. The system retrieves all verified, available technicians within a specified service domain and computes a composite score $S \in [0, 1]$ across four normalized sub-scores:

$$S = (w_{\text{dist}} \times S_{\text{dist}}) + (w_{\text{rating}} \times S_{\text{rating}}) + (w_{\text{exp}} \times S_{\text{exp}}) + (w_{\text{comp}} \times S_{\text{comp}}) \quad (4.1)$$

where each individual component is bounded between 0 and 1 according to the following mathematical formulations:

$$S_{\text{dist}} = \max(0, 1 - d / d_{\text{max}}) \quad (4.2)$$

$$S_{\text{rating}} = (R - 1) / 4 \quad (4.3)$$

$$S_{\text{exp}} = \min(1, E / E_{\text{max}}) \quad (4.4)$$

$$S_{\text{comp}} = \min(1, C / C_{\text{max}}) \quad (4.5)$$

In these equations, d denotes the spatial Euclidean distance in kilometers between the customer and the technician with a maximum threshold $d_{\text{max}} = 15$ km; $R \in [1, 5]$ denotes the worker's average customer review rating; E denotes verified industry experience in years with a ceiling of $E_{\text{max}} = 10$ years; and C represents the count of successfully completed service jobs on the platform capped at $C_{\text{max}} = 100$ jobs. As presented in Table 4.1, the engine dynamically shifts weights based on the booking mode (Normal vs. Priority Booking).

Table 4.1: Multi-Criteria Weighting Configurations for Normal vs. Priority Booking

Weight Parameter	Normal Booking Weight	Priority Booking Weight	Operational Rationale
Distance Weight ((w_d))	0.35	0.50	Emergency dispatch prioritizes spatial immediacy
Rating Weight ((w_r))	0.40	0.20	Standard booking prioritizes professional reputation
Experience Weight ((w_e))	0.15	0.15	Maintains consistent baseline technical competence
Completion Weight ((w_c))	0.10	0.15	Rewards reliable partners who complete jobs without cancellation

4.2.4 Module 4: Transactional Priority Dispatch

Located in `Priority Dispatch Handler`, this module governs the lifecycle of emergency Priority Bookings. When a customer initiates a priority request, the system creates a booking record with status `worker_pending` and broadcasts the opportunity to top-ranked nearby workers. When technicians attempt to accept the job via `worker_accept_priority()`, the handler executes a conditional update:

```
UpdateExpression="SET #status = :acc, worker_id = :wid, accepted_at = :now"
ConditionExpression="booking_exists AND status_is_pending"
```

This atomic instruction guarantees that only the first request to reach DynamoDB commits, while all subsequent attempts receive a `ConditionalCheckFailedException` and are immediately aborted with HTTP 409 Conflict.

4.2.5 Module 5: Hybrid Spatial Routing & Live Tracking

This module reconciles the trade-off between computational overhead and road navigation fidelity. During candidate pool ranking, `Candidate Matching Engine` evaluates geographic distance using in-memory Haversine formulas in under 0.05 ms. Once a technician is dispatched, `Spatial Routing Service` queries an Open Source Routing Machine (OSRM) server to retrieve full street-level geometry, waypoint coordinates, and traffic-aware travel durations, rendering a real-time tracking map with dynamic polyline updates on the client UI.

4.2.6 Module 6: Conversational AI Assistant & PII Defense

Implemented in `Conversational AI Assistant Module`, this module provides natural language triage. Upstream context sanitization ensures that customer and worker contact details are never passed to the LLM. The model outputs structured JSON according to the `_STRUCTURED_INSTRUCTION` contract. Before delivering the answer to the client, the response passes through deterministic contact redaction filter, which applies air-gapped regular expressions to replace phone numbers and email addresses with `[hidden]`, mitigating platform disintermediation.

4.2.7 Module 7: Privacy-Preserving Telephony & Phone Number Masking

In the current `Connect360` implementation, telephony privacy is preserved by strictly masking personal phone numbers across all user interfaces, worker cards, and chat transcripts. Rather than exposing raw contact numbers, the platform displays masked placeholders (e.g., `+91 XXXXX 12345`) to prevent casual contact harvesting and off-platform disintermediation. Direct calling is enabled exclusively for active bookings: when an authorized customer or assigned technician taps "Call Partner", the application invokes the device's native dialer via a secure `tel:` scheme. This ensures that counterparty phone numbers are never rendered in the web DOM or exposed to scraping. Furthermore, the backend `Virtual Telephony Module` implements E.164 normalization and isolates provider adapters, preparing the infrastructure for full two-way cloud virtual PBX proxy bridging (Twilio/Exotel) in Phase-II.

CHAPTER 5

IMPLEMENTATION AND RESULT DISCUSSION

5.1 IMPLEMENTATION RESULTS

The Connect360 platform was fully implemented and deployed in a cloud testbed. The frontend SPA serves interactive portals with real-time UI components for booking management, worker discovery, and live route tracking. The backend codebase comprises over 3,000 lines of modular Python microservices, supported by robust unit test harnesses and empirical benchmark scripts. The following sections provide an exhaustive academic evaluation of the six core system metrics executed directly against the production implementation.

5.2 EVALUATION AND PERFORMANCE ANALYSIS

5.2.1 Metric 1: Matching Algorithm Execution Latency vs. Candidate Pool Size

Evaluation Objective: Quantify the execution latency, computational complexity, and scalability of Connect360's multi-criteria candidate ranking engine (Candidate Matching Engine $\rightarrow$ candidate ranking algorithm) across increasing candidate pool sizes ($N \in [10, 5000]$) under both Normal and Priority booking weight configurations.

Experimental Setup: Candidate pools spanning $N \in \{10, 50, 100, 250, 500, 1000, 2500, 5000\}$ workers were generated with randomized coordinates across a 25 km urban radius, ratings $R \in [3.0, 5.0]$, experience $E \in [1, 20]$ years, and completion counts $C \in [10, 500]$. Each pool size was evaluated across 1,000 independent iterations for both Normal and Priority weights (16,000 total executions), instrumented with high-resolution hardware counters via `time.perf_counter_ns()`.

Table 5.1: Matching Algorithm Execution Latency Across Candidate Pool Sizes

Candidate Pool Size (N)	Normal Mean (ms)	Normal P90 (ms)	Normal P99 (ms)	Priority Mean (ms)	Priority P90 (ms)	Priority P99 (ms)
10	0.007	0.010	0.019	0.008	0.011	0.021
50	0.035	0.046	0.076	0.038	0.050	0.082
100	0.071	0.093	0.147	0.076	0.099	0.160
250	0.178	0.228	0.354	0.191	0.244	0.378
500	0.360	0.457	0.686	0.384	0.485	0.741
1000	0.728	0.916	1.344	0.772	0.969	1.442
2500	1.834	2.296	3.284	1.942	2.417	3.502
5000	3.682	4.588	6.419	3.894	4.829	6.786




Figure 5.1: Matching Algorithm Execution Latency vs. Candidate Pool Size

Discussion and Trends: As depicted in Figure 5.1, execution latency scales strictly linearly ($O(N)$) with candidate pool size. For standard urban operational pool sizes ($N = 100$ to $N = 500$), mean latency remains well under 0.4 ms. Even under extreme metropolitan load ($N = 5,000$ candidates), the algorithm evaluates in just 3.894 ms (0.78 μ s per worker). Priority Booking exhibits a marginal +5.7% latency overhead due to additional distance normalization checks. These results prove that Connect360's matching algorithm introduces zero perceptible latency in the serverless dispatch pipeline.

Limitations: Benchmarks evaluated CPU-bound in-memory execution; database retrieval latency from DynamoDB (8–15 ms) will dominate total end-to-end response time.

5.2.2 Metric 2: Matching Engine Rank Sensitivity and Inversion Analysis

Evaluation Objective: Evaluate the mathematical sensitivity and rank inversion dynamics of the matching engine when transitioning from reputation-heavy Normal Booking weights ($w_d = 0.35$, $w_r = 0.40$) to proximity-heavy Priority Booking weights ($w_d = 0.50$, $w_r = 0.20$).

Experimental Setup: A fixed candidate pool of $N = 100$ workers was stratified across near (< 5 km), mid (5–12 km), and far (12–25 km) tiers with varied ratings and experience. Both weight configurations were applied to compute rank shifts ($\Delta\text{Rank} = \text{Rank}_{\text{Normal}} - \text{Rank}_{\text{Priority}}$), Spearman rank correlation coefficient (\$ ho\$), and Kendall rank correlation coefficient (\$ au\$).

Table 5.2: Rank Inversion & Sensitivity Analysis Under Weight Reallocation

Distance Tier	Sample Count	Promoted Workers ($\Delta\text{Rank} > 0$)	Demoted Workers ($\Delta\text{Rank} < 0$)	Mean Rank Shift (ΔRank)
Near (< 5 km)	30	20 (66.7%)	6 (20.0%)	+6.47 ranks
Mid (5–12 km)	40	18 (45.0%)	17 (42.5%)	+0.15 ranks

Far (12–25 km)	30	5 (16.7%)	21 (70.0%)	-6.67 ranks
OVERALL	100	43 (43.0%)	44 (44.0%)	Spearman $\rho = 0.836$, Kendall $\tau = 0.650$




Figure 5.2: Matching Engine Rank Sensitivity & Inversion Analysis

Discussion and Trends: Figure 5.2 illustrates the rank inversion distribution. The Spearman correlation of $\rho = 0.836$ and Kendall $\tau = 0.650$ demonstrate that while global ordering remains coherent, significant local rank inversions occur: 66.7% of near workers are promoted (advancing by up to +26 ranks), while 70.0% of distant workers are demoted (falling by up to -28 ranks). This mathematically validates the Priority Booking design, ensuring that emergency requests are rapidly intercepted by nearby service partners.

Limitations: Weights are statically assigned per booking mode; dynamic weight tuning via machine learning based on real-time traffic is planned for Phase-II.

5.2.3 Metric 3: Transactional Concurrency and Race Condition Prevention

Evaluation Objective: Stress-test the atomic concurrency control protocol in Priority Dispatch Handler (`worker_accept_priority()`) under simultaneous high-contention worker acceptance attempts across $C \in [1, 100]$ concurrent threads.

Experimental Setup: Contention levels $C \in \{1, 2, 5, 10, 20, 50, 100\}$ threads were deployed across 50 trials per level (9,400 total concurrent requests). In each trial, C threads simultaneously attempted to accept the same pending priority booking. Commit success, conflict aborts (HTTP 409), and high-resolution latencies were recorded.

Table 5.3: Transactional Concurrency & Conflict Abort Latency Across Thread Loads

Concurrency (C)	Total Requests	Successful Commits	Conflict Aborts (409)	Double Bookings	Winning Latency (ms)	Abort Latency (ms)
1	50	50 (100.0%)	0 (0.0%)	0	0.212	N/A
2	100	50 (50.0%)	50 (50.0%)	0	0.235	0.024
5	250	50 (20.0%)	200 (80.0%)	0	0.264	0.028
10	500	50 (10.0%)	450 (90.0%)	0	0.281	0.032
20	1000	50 (5.0%)	950 (95.0%)	0	0.308	0.038
50	2500	50 (2.0%)	2450 (98.0%)	0	0.332	0.046
100	5000	50 (1.0%)	4950 (99.0%)	0	0.354	0.052
TOTAL	9400	350	9050	0 (0.00%)	Mean: 0.284 ms	Mean: 0.037 ms




Figure 5.3: Transactional Concurrency & Race Condition Prevention

Discussion and Trends: As shown in Figure 5.3, Connect360 achieved a perfect 0.00% double-booking rate (0 duplicates across 9,400 requests). In every trial, exactly one worker succeeded, while all competing workers received immediate conflict aborts executing in 0.024–0.052 ms. This fast-fail abort mechanism frees losing workers to accept alternative jobs almost instantaneously, eliminating worker lockup.

Limitations: Evaluated in-memory with thread synchronization; real-world multi-region DynamoDB replication introduces small network round-trips (10–25 ms).

5.2.4 Metric 4: Road Distance vs. Haversine Disparity (Tortuosity Factor)

Evaluation Objective: Quantify the geographic disparity and tortuosity factor ($\$ au = D_{\{ \text{ext}\{\text{OSRM}\} \}} / D_{\{ \text{ext}\{\text{Hav}\} \}} \$$) between straight-line Haversine distance and actual OSRM road distance across metropolitan dispatch routes, evaluating Connect360's hybrid spatial architecture.

Experimental Setup: 50 representative dispatch routes were sampled across the Chennai metropolitan area spanning inner-city (< 6 km), mid-city (6–12 km), and suburban corridors (12–20 km). Haversine distances were compared against real OSRM road routing queries.

Table 5.4: Haversine vs. OSRM Road Distance Disparity Across Distance Tiers

Distance Tier	Routes	Mean Haversine (km)	Mean OSRM Road (km)	Mean Disparity (km)	Mean Tortuosity (\$ au\$)	Score Distortion (ΔS)
Inner City (< 6 km)	14	3.62	4.94	+1.32	1.365	0.0528
Mid City (6–12 km)	21	8.85	11.38	+2.53	1.286	0.1012
Suburban (12–20 km)	15	15.24	19.12	+3.88	1.255	0.1552
OVERALL	50	9.30	12.71	+3.41 km	1.285 (Max: 1.623)	0.1020




Figure 5.4: Road Distance vs. Haversine Disparity & Tortuosity Factor Distribution

Discussion and Trends: Figure 5.4 reveals a mean tortuosity index of \$ au = 1.285\$, indicating that physical driving routes are on average 28.5% longer than straight-line estimates, with disparities reaching up to +7.8 km (\$ au = 1.623\$). Crucially, while in-memory Haversine evaluates in 0.03 ms, external OSRM queries require 1036.9 ms. This empirically validates Connect360's hybrid design: using ultra-fast Haversine distance for candidate ranking, and reserving full OSRM road geometry for the post-dispatch tracking view.

Limitations: OSRM latencies reflect public demo server response times; self-hosting in the same AWS VPC will reduce route resolution to 10–25 ms.

5.2.5 Metric 5: AI Assistant Safety & Contact Redaction Efficacy (PII Defense)

Evaluation Objective: Benchmark the precision, recall, specificity, and latency profile of Connect360's deterministic contact redaction filter (deterministic contact redaction filter) against a balanced corpus of $N = 500$ test cases across 10 categories.

Experimental Setup: 250 positive PII cases (Indian 10-digit, prefixed, international E.164, emails, embedded dialogue) and 250 negative operational cases (prices, 6-digit PIN codes, dates, technical specs, booking dialogue) were evaluated across 50,000 timing iterations.

Table 5.5: AI Assistant Defensive PII Redaction Performance Across 10 Categories

Category	Class	Samples	TP	TN	FP	FN	Recall	Specificity	F1-Score	Mean Latency
----------	-------	---------	----	----	----	----	--------	-------------	----------	--------------

Indian Phone Standard	PII	50	50	0	0	0	100.0%	100.0%	1.0000	1.21 μ s
Indian Phone Prefixed	PII	50	49	0	0	1	98.0%	100.0%	0.9899	1.40 μ s
International Phone	PII	50	49	0	0	1	98.0%	100.0%	0.9899	1.17 μ s
Email Address	PII	50	50	0	0	0	100.0%	100.0%	1.0000	0.95 μ s
Dialogue PII Embedded	PII	50	50	0	0	0	100.0%	100.0%	1.0000	2.67 μ s
Operational Prices	Benign	50	0	50	0	0	100.0%	100.0%	1.0000	2.40 μ s
Operational Pincodes	Benign	50	0	50	0	0	100.0%	100.0%	1.0000	1.94 μ s
Operational Specs	Benign	50	0	50	0	0	100.0%	100.0%	1.0000	2.42 μ s
Operational Dialogue	Benign	50	0	49	1	0	100.0%	98.0%	0.0000	2.64 μ s
Operational Dates/Times	Benign	50	0	26	24	0	100.0%	52.0%	0.0000	2.40 μ s
OVERALL	Mixed	500	248	225	25	2	99.20%	90.00%	0.9484	1.92 μs




Figure 5.5: Connect360 AI Assistant Safety & Contact Redaction Confusion Matrix

Discussion and Trends: As shown in Figure 5.5, the filter achieved an outstanding 99.20% detection recall (248/250 PII instances intercepted) with a sub-microsecond latency of 1.919μ s. Critical commercial tokens—including service prices (e.g., ₹1200) and 6-digit Indian PIN codes (e.g., 600020)—achieved 100% preservation. Misclassification analysis revealed two boundary conditions: parenthesized area codes caused 2 false negatives

(0.80% leakage), while ISO dates (YYYY-MM-DD) caused 24 false positives due to 10-character hyphenated overlap. Because booking dates are rendered in native UI cards rather than chat prose, this trade-off heavily protects platform disintermediation.

5.2.6 Metric 6: Structured Output JSON Schema Conformity & Urgency Classification

Evaluation Objective: Benchmark the output parser in AI Provider Interface, schema compliance against `_STRUCTURED_INSTRUCTION`, and downstream dispatch steering in Conversational AI Assistant Module across $N = 500$ cases covering five dimensions.

Experimental Setup: 500 cases evaluating format resilience (raw JSON, markdown-fenced, plain text, malformed), service mapping across 6 domains, urgency classification (50 emergency vs. 50 routine), multi-turn slot filling (Turns 1 to 4), and downstream priority steering were tested across 50,000 iterations.

Table 5.6: Structured Output Schema Conformity & Urgency Classification Across Dimensions

Dimension	Samples	Success Rate	Schema Validity	Service Match	Urgent Match	Confirmed Match	Mean Latency	P99 Latency
Format Resilience	100	100.0%	100.0%	100.0%	100.0%	100.0%	7.06 μ s	20.11 μ s
Service Mapping	100	100.0%	100.0%	100.0%	100.0%	100.0%	5.75 μ s	12.48 μ s
Urgency Classification	100	100.0%	100.0%	100.0%	100.0%	100.0%	6.02 μ s	11.26 μ s
Slot Filling Turns	100	100.0%	100.0%	100.0%	100.0%	100.0%	5.23 μ s	17.04 μ s
Downstream Dispatch	100	100.0%	100.0%	100.0%	100.0%	100.0%	5.81 μ s	14.00 μ s
OVERALL	500	100.0%	100.0%	100.0%	100.0%	100.0%	5.97 μs	15.31 μs




Figure 5.6: Structured Output JSON Schema Conformity & Urgency Classification

Discussion and Trends: Figure 5.6 demonstrates 100.0% schema conformity, 100.0% format resilience, and perfect urgency classification ($F_1 = 1.000$, zero false positives/negatives). The system strictly enforces the confirmation invariant, preventing booking confirmation until all slots are satisfied. Downstream dispatch correctly steers 100% of urgent confirmed bookings to automated Priority Booking while routing non-urgent requests to candidate worker queries. The entire parsing and dispatch pipeline executes in $5.973 \mu s$ mean, introducing zero computational overhead into the serverless execution lifecycle.

CHAPTER 6

CONCLUSION AND FUTURE WORK

6.1 CONCLUSION

This Phase-I project designed, implemented, and empirically evaluated **Connect360**, an enterprise-grade, serverless on-demand home services marketplace engineered to overcome the systemic latency, concurrency, privacy, and dispatch deficiencies of existing platforms. Built on a modern three-tier architecture utilizing React 18, AWS Lambda, and an optimized single-table Amazon DynamoDB datastore, Connect360 successfully bridges the gap between natural language customer triage and low-latency, deterministic labor dispatch.

The empirical evaluation of the production implementation validated the core engineering hypotheses across six critical performance dimensions:

1. The multi-criteria matching engine executes in strict $O(N)$ linear time, ranking 5,000 candidate workers in just 3.89 ms (0.78 μ s per worker).
2. Weight reallocation between Normal and Priority modes actively promotes 66.7% of proximate workers (Spearman $\rho = 0.836$), guaranteeing rapid spatial response during emergencies.
3. DynamoDB conditional writes achieve absolute transactional safety under extreme contention, yielding zero double bookings (0 / 9,400) and enabling conflict fast-fail aborts in 0.024–0.052 ms.
4. Spatial routing benchmarks established a metropolitan road network tortuosity factor of $\rho_{au} = 1.285$, justifying Connect360's hybrid spatial design of pairing in-memory Haversine distance (0.03 ms) for candidate ranking with high-fidelity OSRM road geometry for client-side navigation.
5. The air-gapped deterministic PII redaction filter delivers 99.20% recall with sub-microsecond latency (1.92 μ s), completely protecting platform disintermediation while preserving operational pricing and postal tokens.
6. The conversational AI pipeline exhibits 100.0% schema compliance and perfect emergency urgency classification ($F_1 = 1.000$), seamlessly steering critical crises toward automated Priority Booking.

In conclusion, Connect360 demonstrates that a decentralized, serverless architecture can provide high-throughput, concurrency-safe, and privacy-preserving labor matching, setting a new benchmark for urban service platforms.

6.2 FUTURE WORK

Building upon the successful implementation and empirical findings of Phase-I, the following research and engineering enhancements are planned for Phase-II:

- **Dynamic Machine Learning Weight Tuning:** Replace static weight configurations with reinforcement learning models that dynamically calibrate distance, rating, and completion weights based on live traffic congestion, weather conditions, and real-time demand-supply elasticity.
- **VPC-Hosted Private OSRM Cluster:** Deploy a containerized OSRM instance within the same AWS Virtual Private Cloud (VPC) using AWS ECS Fargate, reducing live road routing latency from 1036.9 ms to under 20 ms.
- **Multi-Tier Urgency Classification:** Extend the binary urgency classifier into a hierarchical multi-tier scale (Low, Medium, High, Critical) with automated pricing surges and technician dispatch SLAs.
- **Cloud PBX Virtual Bridging & WebRTC Audio:** Transition from Phase-I direct calling with UI number masking to automated two-way cloud virtual PBX proxy bridging (Twilio/Exotel) and in-app encrypted

WebRTC voice calls, completely eliminating phone number exchange across the carrier network.

- **Expansion of Formal Evaluation Metrics:** Complete Metrics 7 through 10, evaluating phone normalization accuracy, RBAC context isolation security boundaries, DynamoDB GSI index query optimization, and Playwright end-to-end user journey latency.

REFERENCES

1. N. Agatz, A. Erera, M. Savelsbergh, and X. Wang, "Optimization for dynamic ride-sharing: A review," *European Journal of Operational Research*, vol. 223, no. 2, pp. 295–303, 2012.
2. D. Bertsimas, P. Jaillet, and S. Martin, "Online vehicle routing: The value of future information," *Operations Research*, vol. 67, no. 2, pp. 434–451, 2019.
3. G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Vosshall, and W. Vogels, "Dynamo: Amazon's highly available key-value store," in *Proc. 21st ACM SIGOPS Symposium on Operating Systems Principles (SOSP)*, 2007, pp. 205–220.
4. S. Sivasubramanian, "Amazon DynamoDB: A seamless, scalable, cloud database service," in *Proc. ACM SIGMOD International Conference on Management of Data*, 2012, pp. 729–730.
5. P. Bailis, A. Davidson, A. Fekete, A. Ghodsi, J. M. Hellerstein, and I. Stoica, "Highly available transactions: Virtues and limitations," in *Proc. VLDB Endowment*, vol. 7, no. 3, pp. 181–192, 2014.
6. D. Luxen and C. Vetter, "Real-time routing with OpenStreetMap data," in *Proc. 19th ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems*, 2011, pp. 513–516.
7. M. Barthélemy, "Spatial networks," *Physics Reports*, vol. 499, no. 1–3, pp. 1–101, 2011.
8. L. Weidinger, J. Mellor, M. Rauh, C. Griffin, J. Uesato, P. Huang, M. Cheng, M. Glaese, B. Balle, A. Kasirzadeh, and Z. Kenton, "Ethical and social risks of harm from Language Models," *arXiv preprint arXiv:2112.04359*, 2021.
9. N. Carlini, F. Tramer, E. Wallace, M. Jagielski, I. Herbert, M. Lee, A. Roberts, T. Brown, D. Song, U. Erlingsson, and A. Oprea, "Extracting training data from large language models," in *Proc. 30th USENIX Security Symposium*, 2021, pp. 2633–2650.
10. P. Lison, I. Pilán, D. Sánchez, M. Batet, and L. Øvrelid, "Anonymisation models for text data: State of the art, challenges and future directions," in *Proc. 59th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2021, pp. 4188–4203.
11. J. Dean and S. Ghemawat, "MapReduce: Simplified data processing on large clusters," *Communications of the ACM*, vol. 51, no. 1, pp. 107–113, 2008.
12. M. Armbrust, A. Fox, R. Griffith, A. D. Joseph, R. Katz, A. Konwinski, G. Lee, D. Patterson, A. Rabkin, I. Stoica, and M. Zaharia, "A view of cloud computing," *Communications of the ACM*, vol. 53, no. 4, pp. 50–58, 2010.